

In [1]:

```
import warnings
warnings.filterwarnings("ignore")
import pandas as pd
import sqlite3
import csv
import matplotlib.pyplot as plt
import seaborn as sns
import numpy as np
from wordcloud import WordCloud
import re
import os
from sqlalchemy import create_engine # database connection
import datetime as dt
from nltk.corpus import stopwords
from nltk.tokenize import word_tokenize
from nltk.stem.snowball import SnowballStemmer
from sklearn.feature_extraction.text import CountVectorizer
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.multiclass import OneVsRestClassifier
from sklearn.linear_model import SGDClassifier
from sklearn import metrics
from sklearn.metrics import f1_score, precision_score, recall_score
from sklearn import svm
from sklearn.linear_model import LogisticRegression

from datetime import datetime
from sklearn.externals import joblib

from sklearn.model_selection import GridSearchCV
```

3.3 Cleaning and preprocessing of Questions

3.3.1 Preprocessing

1. Sample 1M data points
2. Separate out code-snippets from Body
3. Remove Special characters from Question title and description (not in code)
4. Remove stop words (Except 'C')
5. Remove HTML Tags
6. Convert all the characters into small letters
7. Use SnowballStemmer to stem the words

In [2]:

```
def striphtml(data):  
    cleanr = re.compile('<.*?>')  
    cleantext = re.sub(cleanr, ' ', str(data))  
    return cleantext  
stop_words = set(stopwords.words('english'))  
stemmer = SnowballStemmer("english")
```

In [3]:

```
#http://www.sqlitetutorial.net/sqlite-python/create-tables/
def create_connection(db_file):
    """ create a database connection to the SQLite database
        specified by db_file
    :param db_file: database file
    :return: Connection object or None
    """
    try:
        conn = sqlite3.connect(db_file)
        return conn
    except Error as e:
        print(e)

    return None

def create_table(conn, create_table_sql):
    """ create a table from the create_table_sql statement
    :param conn: Connection object
    :param create_table_sql: a CREATE TABLE statement
    :return:
    """
    try:
        c = conn.cursor()
        c.execute(create_table_sql)
    except Error as e:
        print(e)

def checkTableExists(dbcon):
    cursr = dbcon.cursor()
    str = "select name from sqlite_master where type='table'"
    table_names = cursr.execute(str)
    print("Tables in the database:")
    tables = table_names.fetchall()
    print(tables[0][0])
    return(len(tables))

def create_database_table(database, query):
    conn = create_connection(database)
    if conn is not None:
        create_table(conn, query)
        checkTableExists(conn)
    else:
        print("Error! cannot create the database connection.")
    conn.close()

sql_create_table = """CREATE TABLE IF NOT EXISTS QuestionsProcessed (question
create_database_table("Processed.db", sql_create_table)
```

Tables in the database:

4. Machine Learning Models

4.1 Converting tags for multilabel problems

X	y1	y2	y3	y4
x1	0	1	1	0
x1	1	0	0	0
x1	0	1	0	0

4.5 Modeling with less data points (0.5M data points) and more weight to title and 500 tags only.

In [4]:

```
sql_create_table = """CREATE TABLE IF NOT EXISTS QuestionsProcessed (question_id INT, title TEXT, tags TEXT, processed BOOLEAN)"""
create_database_table("Titlemoreweight.db", sql_create_table)
```

Tables in the database:
QuestionsProcessed

In [5]:

```
# http://www.sqlitetutorial.net/sqlite-delete/
# https://stackoverflow.com/questions/2279706/select-random-row-from-a-sqlite

read_db = 'train_no_dup.db'
write_db = 'Titlemoreweight.db'
train_datasize = 400000
if os.path.isfile(read_db):
    conn_r = create_connection(read_db)
    if conn_r is not None:
        reader = conn_r.cursor()
        # for selecting first 0.5M rows
        reader.execute("SELECT Title, Body, Tags From no_dup_train LIMIT 5000")
        # for selecting random points
        #reader.execute("SELECT Title, Body, Tags From no_dup_train ORDER BY")

if os.path.isfile(write_db):
    conn_w = create_connection(write_db)
    if conn_w is not None:
        tables = checkTableExists(conn_w)
        writer = conn_w.cursor()
        if tables != 0:
            writer.execute("DELETE FROM QuestionsProcessed WHERE 1")
            print("Cleared All the rows")
```

Tables in the database:
QuestionsProcessed
Cleared All the rows

4.5.1 Preprocessing of questions

1. Separate Code from Body
2. Remove Special characters from Question title and description (not in code)
3. **Give more weightage to title : Add title three times to the question**

- Remove stop words (Except 'C')
- Remove HTML Tags
- Convert all the characters into small letters
- Use SnowballStemmer to stem the words

In [6]:

```
#http://www.bernzilla.com/2008/05/13/selecting-a-random-row-from-an-sqlite-to
start = datetime.now()
preprocessed_data_list=[]
reader.fetchone()
questions_with_code=0
len_pre=0
len_post=0
questions_proccesed = 0
for row in reader:

    is_code = 0

    title, question, tags = row[0], row[1], str(row[2])

    if '<code>' in question:
        questions_with_code+=1
        is_code = 1
    x = len(question)+len(title)
    len_pre+=x

    code = str(re.findall(r'<code>(.*?)</code>', question, flags=re.DOTALL))

    question=re.sub('<code>(.*?)</code>', '', question, flags=re.MULTILINE|re
    question=striphtml(question.encode('utf-8'))

    title=title.encode('utf-8')

    # adding title three time to the data to increase its weight
    # add tags string to the training data

    question=str(title)+" "+str(title)+" "+str(title)+" "+question

    # if questions_proccesed<=train_datasize:
    #     question=str(title)+" "+str(title)+" "+str(title)+" "+question+" "+
    # else:
    #     question=str(title)+" "+str(title)+" "+str(title)+" "+question

    question=re.sub(r'^A-Za-z0-9#+.\-]+', ' ',question)
    words=word_tokenize(str(question.lower()))

    #Removing all single letter and and stopwords from question exceptt for t
    question=' '.join(str(stemmer.stem(j)) for j in words if j not in stop_w

    len_post+=len(question)
    tup = (question,code,tags,x,len(question),is_code)
    questions_proccesed += 1
    writer.execute("insert into QuestionsProcessed(question,code,tags,words_p
    if (questions_proccesed%100000==0):
        print("number of questions completed=",questions_proccesed)
```

```

no_dup_avg_len_pre=(len_pre*1.0)/questions_proccesed
no_dup_avg_len_post=(len_post*1.0)/questions_proccesed

print( "Avg. length of questions(Title+Body) before processing: %d"%no_dup_avg_len_pre)
print( "Avg. length of questions(Title+Body) after processing: %d"%no_dup_avg_len_post)
print( "Percent of questions containing code: %d"%((questions_with_code*100.0)/questions_proccesed))

print("Time taken to run this cell :", datetime.now() - start)

```

```

number of questions completed= 100000
number of questions completed= 200000
number of questions completed= 300000
number of questions completed= 400000
number of questions completed= 500000
Avg. length of questions(Title+Body) before processing: 1239
Avg. length of questions(Title+Body) after processing: 424
Percent of questions containing code: 57
Time taken to run this cell : 0:14:13.702826

```

In [7]:

```

# never forget to close the conections or else we will end up with database l
conn_r.commit()
conn_w.commit()
conn_r.close()
conn_w.close()

```

___ Sample quesitons after preprocessing of data ___

In [8]:

```
if os.path.isfile(write_db):
    conn_r = create_connection(write_db)
    if conn_r is not None:
        reader = conn_r.cursor()
        reader.execute("SELECT question From QuestionsProcessed LIMIT 10")
        print("Questions after preprocessed")
        print('='*100)
        reader.fetchone()
        for row in reader:
            print(row)
            print('-'*100)
conn_r.commit()
conn_r.close()
```

Questions after preprocessed

```
=====
=====
('dynam datagrid bind silverlight dynam datagrid bind silverli
ght dynam datagrid bind silverlight bind datagrid dynam code w
rote code debug code block seem bind correct grid come column
form come grid column although necessari bind nthank repli adv
ance..',)
-----
-----
('java.lang.noclassdeffounderror javax servlet jsp tagext tagl
ibraryvalid java.lang.noclassdeffounderror javax servlet jsp t
agext taglibraryvalid java.lang.noclassdeffounderror javax ser
vlet jsp tagext taglibraryvalid follow guid link instal jstl g
ot follow error tri launch jsp page java.lang.noclassdeffounde
rror javax servlet jsp tagext taglibraryvalid taglib declar in
stal jstl 1.1 tomcat webapp tri project work also tri version
1.2 jstl still messag caus solv',)
-----
-----
('java.sql.sqlexcept microsoft odbc driver manag invalid descr
iptor index java.sql.sqlexcept microsoft odbc driver manag inv
alid descriptor index java.sql.sqlexcept microsoft odbc driver
manag invalid descriptor index use follow code display caus so
lv',)
-----
-----
('better way updat feed fb php sdk better way updat feed fb ph
p sdk better way updat feed fb php sdk novic facebook api read
mani tutori still confused.i find post feed api method like co
rrect second way use curl someth like way better',)
-----
-----
('btnadd click event open two window record ad btnadd click ev
ent open two window record ad btnadd click event open two wind
ow record ad open window search.aspx use code hav add button s
```



```
earch.aspx nwhen insert record btnadd click event open anoth w  
indow nafter insert record close window',)
```

```
-----  
-----  
( 'sql inject issu prevent correct form submiss php sql inject  
issu prevent correct form submiss php sql inject issu prevent  
correct form submiss php check everyth think make sure input f  
ield safe type sql inject good news safe bad news one tag mess  
form submiss place even touch life figur exact html use templa  
t file forgiv okay entir php script get execut see data post n  
one forum field post problem use someth titl field none data g  
et post current use print post see submit noth work flawless s  
tatement though also mention script work flawless local machin  
use host come across problem state list input test mess',)
```

```
-----  
-----  
( 'countabl subaddit lebesgu measur countabl subaddit lebesgu m  
easur countabl subaddit lebesgu measur let lbrace rbrace seque  
nc set sigma -algebra mathcal want show left bigcup right leq  
sum left right countabl addit measur defin set sigma algebra m  
athcal think use monoton properti somewher proof start appreci  
littl helpnthank ad han answer make follow addit construct gi  
ven han answer clear bigcup bigcup cap emptyset neq left bigcu  
p right left bigcup right sum left right also construct subset  
monoton left right leq left right final would sum leq sum resu  
lt follow',)
```

```
-----  
-----  
( 'hql equival sql queri hql equival sql queri hql equival sql  
queri hql queri replac name class properti name error occur hq  
l error',)
```

```
-----  
-----  
( 'undefin symbol architectur i386 objc class skpsmtpmessag ref  
erenc error undefin symbol architectur i386 objc class skpsmt  
messag referenc error undefin symbol architectur i386 objc cla  
ss skpsmtpmessag referenc error import framework send email ap  
plic background import framework i.e skpsmtpmessag somebodi su  
ggest get error collect2 ld return exit status import framewor  
k correct sorc taken framework follow mfmailcomposeviewcontrol  
question lock field updat answer drag drop folder project clic  
k copi nthat',)
```

__ Saving Preprocessed data to a Database __

In [9]:

```
#Taking 0.5 Million entries to a dataframe.
write_db = 'Titlmoreweight.db'
if os.path.isfile(write_db):
    conn_r = create_connection(write_db)
    if conn_r is not None:
        preprocessed_data = pd.read_sql_query("""SELECT question, Tags FROM C
conn_r.commit()
conn_r.close()
```

In [10]:

```
preprocessed_data.head()
```

Out[10]:

	question	tags
0	dynam datagrid bind silverlight dynam datagrid...	c# silverlight data-binding
1	dynam datagrid bind silverlight dynam datagrid...	c# silverlight data-binding columns
2	java.lang.noclassdeffounderror javax servlet j...	jsp jstl
3	java.sql.sqlexcept microsoft odbc driver manag...	java jdbc
4	better way updat feed fb php sdk better way up...	facebook api facebook-php-sdk

In [11]:

```
print("number of data points in sample :", preprocessed_data.shape[0])
print("number of dimensions :", preprocessed_data.shape[1])
```

```
number of data points in sample : 500000
number of dimensions : 2
```

__ Converting string Tags to multilable output variables __

In [12]:

```
vectorizer = CountVectorizer(tokenizer = lambda x: x.split(), binary='true')
multilabel_y = vectorizer.fit_transform(preprocessed_data['tags'])
```

__ Selecting 500 Tags __

In [13]:

```
def tags_to_choose(n):
    t = multilabel_y.sum(axis=0).tolist()[0]
    sorted_tags_i = sorted(range(len(t)), key=lambda i: t[i], reverse=True)
    multilabel_yn=multilabel_y[:,sorted_tags_i[:n]]
    return multilabel_yn


def questions_explained_fn(n):
    multilabel_yn = tags_to_choose(n)
    x= multilabel_yn.sum(axis=1)
    return (np.count_nonzero(x!=0))
```

In [14]:

```
questions_explained = []
total_tags=multilabel_y.shape[1]
total_qs=preprocessed_data.shape[0]
for i in range(500, total_tags, 100):
    questions_explained.append(np.round(((total_qs-questions_explained_fn(i))
```

In [15]:

```
fig, ax = plt.subplots()
ax.plot(questions_explained)
xlabel = list(500+np.array(range(-50,450,50))*50)
ax.set_xticklabels(xlabel)
plt.xlabel("Number of tags")
plt.ylabel("Number Questions covered partially")
plt.grid()
plt.show()
# you can choose any number of tags based on your computing power, minimum is
print("with ",5500,"tags we are covering ",questions_explained[50],"% of ques
print("with ",500,"tags we are covering ",questions_explained[0],"% of ques
```



with 5500 tags we are covering 99.157 % of questions
with 500 tags we are covering 90.956 % of questions

In [16]:

```
# we will be taking 500 tags
multilabel_yx = tags_to_choose(500)
print("number of questions that are not covered :", questions_explained_fn(500))
```

number of questions that are not covered : 45221 out of 50000
0

In [17]:

```
x_train=preprocessed_data.head(train_datasize)
x_test=preprocessed_data.tail(preprocessed_data.shape[0] - 400000)

y_train = multilabel_yx[0:train_datasize,:]
y_test = multilabel_yx[train_datasize:preprocessed_data.shape[0],:]
```

In [18]:

```
print("Number of data points in train data :", y_train.shape)
print("Number of data points in test data :", y_test.shape)
```

Number of data points in train data : (400000, 500)

Number of data points in test data : (100000, 500)

4.5.2 Featurizing data with Tfidf vectorizer

In [23]:

```
start = datetime.now()
vectorizer = TfidfVectorizer(min_df=0.00009, max_features=200000, smooth_idf=1,
                             tokenizer = lambda x: x.split(), sublinear_tf=True)
x_train_multilabel = vectorizer.fit_transform(x_train['question'])
x_test_multilabel = vectorizer.transform(x_test['question'])
print("Time taken to run this cell :", datetime.now() - start)
```

Time taken to run this cell : 0:02:51.543164

In [24]:

```
print("Dimensions of train data X:", x_train_multilabel.shape, "Y :", y_train.shape)
print("Dimensions of test data X:", x_test_multilabel.shape, "Y:", y_test.shape)
```

Dimensions of train data X: (400000, 94928) Y : (400000, 500)

Dimensions of test data X: (99998, 94928) Y: (99998, 500)

4.5.3 Applying Logistic Regression with OneVsRest Classifier

In [25]:

```
start = datetime.now()
classifier = OneVsRestClassifier(SGDClassifier(loss='log', alpha=0.00001, per
classifier.fit(x_train_multilabel, y_train)
predictions = classifier.predict (x_test_multilabel)

print("Accuracy :",metrics.accuracy_score(y_test, predictions))
print("Hamming loss ",metrics.hamming_loss(y_test,predictions))

precision = precision_score(y_test, predictions, average='micro')
recall = recall_score(y_test, predictions, average='micro')
f1 = f1_score(y_test, predictions, average='micro')

print("Micro-average quality numbers")
print("Precision: {:.4f}, Recall: {:.4f}, F1-measure: {:.4f}".format(precisio

precision = precision_score(y_test, predictions, average='macro')
recall = recall_score(y_test, predictions, average='macro')
f1 = f1_score(y_test, predictions, average='macro')

print("Macro-average quality numbers")
print("Precision: {:.4f}, Recall: {:.4f}, F1-measure: {:.4f}".format(precisio

print (metrics.classification_report(y_test, predictions))
print("Time taken to run this cell :", datetime.now() - start)
```

Accuracy : 0.23705474109482189

Hamming loss 0.0027796155923118463

Micro-average quality numbers

Precision: 0.7220, Recall: 0.3258, F1-measure: 0.4490

Macro-average quality numbers

Precision: 0.5483, Recall: 0.2582, F1-measure: 0.3350

	precision	recall	f1-score	support
0	0.94	0.64	0.76	5519
1	0.69	0.26	0.37	8189
2	0.81	0.37	0.51	6529
3	0.81	0.43	0.56	3231
4	0.81	0.41	0.54	6430
5	0.81	0.34	0.48	2878
6	0.87	0.49	0.63	5086
7	0.88	0.54	0.67	4533
8	0.60	0.13	0.21	3000
9	0.81	0.53	0.64	2765
10	0.59	0.16	0.25	3051
11	0.70	0.22	0.45	3000

In [28]:

```
joblib.dump(classifier, 'lr_with_more_title_weight.pkl')
```

Out[28]:

```
['lr_with_more_title_weight.pkl']
```

In [29]:

```
start = datetime.now()
classifier_2 = OneVsRestClassifier(LogisticRegression(penalty='l1'), n_jobs=-1)
classifier_2.fit(x_train_multilabel, y_train)
predictions_2 = classifier_2.predict(x_test_multilabel)
print("Accuracy :", metrics.accuracy_score(y_test, predictions_2))
print("Hamming loss ", metrics.hamming_loss(y_test, predictions_2))

precision = precision_score(y_test, predictions_2, average='micro')
recall = recall_score(y_test, predictions_2, average='micro')
f1 = f1_score(y_test, predictions_2, average='micro')

print("Micro-average quality numbers")
print("Precision: {:.4f}, Recall: {:.4f}, F1-measure: {:.4f}".format(precision, recall, f1))

precision = precision_score(y_test, predictions_2, average='macro')
recall = recall_score(y_test, predictions_2, average='macro')
f1 = f1_score(y_test, predictions_2, average='macro')

print("Macro-average quality numbers")
print("Precision: {:.4f}, Recall: {:.4f}, F1-measure: {:.4f}".format(precision, recall, f1))

print(metrics.classification_report(y_test, predictions_2))
print("Time taken to run this cell :", datetime.now() - start)
```

```
Accuracy : 0.251065021300426
Hamming loss  0.0027028740574811497
Micro-average quality numbers
Precision: 0.7172, Recall: 0.3673, F1-measure: 0.4858
Macro-average quality numbers
Precision: 0.5571, Recall: 0.2951, F1-measure: 0.3711
```

	precision	recall	f1-score	support
0	0.94	0.72	0.82	5519
1	0.70	0.34	0.45	8189
2	0.80	0.42	0.55	6529
3	0.82	0.49	0.61	3231
4	0.80	0.44	0.57	6430
5	0.82	0.38	0.52	2878
6	0.86	0.53	0.66	5086
7	0.87	0.58	0.70	4533
8	0.60	0.13	0.22	3000
9	0.82	0.57	0.67	2765
10	0.60	0.20	0.30	3051
11	0.60	0.20	0.30	3051

1. Use bag of words upto 4 grams

In [19]:

```
vectorizer_bow = CountVectorizer(min_df=0.00009, ngram_range=(1,4), max_featu
vectorizer_bow.fit(x_train['question'].values) # fit has to happen only on tr

X_train_bow = vectorizer_bow.transform(x_train['question'].values)
X_test_bow = vectorizer_bow.transform(x_test['question'].values)

print("After vectorizations")
print(X_train_bow.shape, y_train.shape)
print(X_test_bow.shape, y_test.shape)
```

After vectorizations
(400000, 99057) (400000, 500)
(100000, 99057) (100000, 500)

In []:

```
#Saving bows
'''from sklearn.externals import joblib
#data points with 0.5 million data
joblib.dump(X_train_bow, 'X_train_bow.pkl')
joblib.dump(X_test_bow, 'X_test_bow.pkl')

#target class i.e multilabel classes with 0.5 million
joblib.dump(y_train, 'y_train.pkl')
joblib.dump(y_test, 'y_test.pkl')'''
```

In [2]:

```
from sklearn.externals import joblib
X_train_bow = joblib.load('X_train_bow.pkl')
X_test_bow = joblib.load('X_test_bow.pkl')
y_train = joblib.load('y_train.pkl')
y_test = joblib.load('y_test.pkl')
```

In [3]:

```
print(X_train_bow.shape, y_train.shape)
print(X_test_bow.shape, y_test.shape)
```

(400000, 99057) (400000, 500)
(100000, 99057) (100000, 500)

2. Perform hyperparameter tuning for Logistic regression using GridSearch


In [20]:

```
start = datetime.now()
parameters = {'estimator__C' : [1e-06, 1e-04, 1e-02, 1, 100]}
ovr = OneVsRestClassifier(LogisticRegression())
clf = GridSearchCV(ovr, parameters, cv=2, scoring='f1_micro', return_train_score=True)
clf.fit(X_train_bow, y_train)

results = pd.DataFrame.from_dict(clf.cv_results_)
results = results.sort_values(['param_estimator__C'])

train_f1_micro= results['mean_train_score']
train_f1_micro_std= results['std_train_score']
cv_f1_micro = results['mean_test_score']
cv_f1_micro_std= results['std_test_score']
K = results['param_estimator__C']

plt.plot(K, train_f1_micro, label='Train F1 Micro')
plt.plot(K, cv_f1_micro, label='CV F1 Micro')
plt.scatter(K, train_f1_micro, label='Train F1 Micro points')
plt.scatter(K, cv_f1_micro, label='CV F1 Micro points')
plt.legend()
plt.xlabel("C : hyperparameter")
plt.xscale('log')
plt.ylabel("F1 Micro")
plt.title("Hyper parameter Vs F1 Micro plot")
plt.grid()
plt.show()
print("Time taken to run this cell :", datetime.now() - start)
```



Time taken to run this cell : 3:33:54.567605

Compute the micro f1 score with Logistic regression(OvR)

In [10]:

```
classifier_2 = OneVsRestClassifier(LogisticRegression(C = 0.01), n_jobs=-1)
classifier_2.fit(X_train_bow, y_train)
predictions_2 = classifier_2.predict(X_test_bow)
print("Accuracy :", metrics.accuracy_score(y_test, predictions_2))
print("Hamming loss ", metrics.hamming_loss(y_test, predictions_2))

precision = precision_score(y_test, predictions_2, average='micro')
recall = recall_score(y_test, predictions_2, average='micro')
f1 = f1_score(y_test, predictions_2, average='micro')

print("Micro-average quality numbers")
print("Precision: {:.4f}, Recall: {:.4f}, F1-measure: {:.4f}".format(precision, recall, f1))

precision = precision_score(y_test, predictions_2, average='macro')
recall = recall_score(y_test, predictions_2, average='macro')
f1 = f1_score(y_test, predictions_2, average='macro')

print("Macro-average quality numbers")
print("Precision: {:.4f}, Recall: {:.4f}, F1-measure: {:.4f}".format(precision, recall, f1))

print(metrics.classification_report(y_test, predictions_2))
```

44	0.60	0.45	0.54	620
45	0.68	0.32	0.43	852
46	0.72	0.35	0.47	534
47	0.38	0.17	0.23	350
48	0.72	0.49	0.58	496
49	0.79	0.57	0.66	785
50	0.23	0.05	0.09	475
51	0.37	0.16	0.22	305
52	0.40	0.02	0.05	251
53	0.65	0.38	0.48	914
54	0.47	0.17	0.25	728
55	0.31	0.03	0.06	258
56	0.44	0.21	0.28	821
57	0.42	0.11	0.17	541
58	0.77	0.26	0.39	748
59	0.94	0.66	0.78	724
60	0.36	0.07	0.12	660
61	0.75	0.21	0.33	235
62	0.91	0.67	0.77	718
63	0.84	0.62	0.72	468

Hyperparameter tuning Linear-SVM (SGDClassifier with loss-hinge)

In [24]:

```
parameters = {'estimator__alpha' : [1e-06, 1e-04, 1e-02, 1, 100]}
ovr = OneVsRestClassifier(SGDClassifier(loss='hinge'))
clf = GridSearchCV(ovr, parameters, cv=2, scoring='f1_micro', return_train_score=True)
clf.fit(X_train_bow, y_train)

results_2 = pd.DataFrame.from_dict(clf.cv_results_)
results_2 = results_2.sort_values(['param_estimator__alpha'])

train_f1_micro= results_2['mean_train_score']
train_f1_micro_std= results_2['std_train_score']
cv_f1_micro = results_2['mean_test_score']
cv_f1_micro_std= results_2['std_test_score']
K = results_2['param_estimator__alpha']

plt.plot(K, train_f1_micro, label='Train F1 Micro')
plt.plot(K, cv_f1_micro, label='CV F1 Micro')
plt.scatter(K, train_f1_micro, label='Train F1 Micro points')
plt.scatter(K, cv_f1_micro, label='CV F1 Micro points')
plt.legend()
plt.xlabel("alpha : hyperparameter")
plt.xscale('log')
plt.ylabel("F1 Micro")
plt.title("Hyper parameter Vs F1 Micro plot")
plt.grid()
plt.show()
```


KeyError Traceback (most recent call last)

<ipython-input-24-3fb0d96df475> in <module>

```
5
6 results_2 = pd.DataFrame.from_dict(clf.cv_results_)
----> 7 results_2 = results_2.sort_values(['param_estimator__alpha'])
8
9 train_f1_micro= results_2['mean_train_score']
```

c:\users\addu\appdata\local\programs\python\python37\lib\site-packages\pandas\core\frame.py in sort_values(self, by, axis, ascending, inplace, kind, na_position)

```
5006
5007         by = by[0]
-> 5008         k = self._get_label_or_level_values(by, axis=axis)
5009
5010         if isinstance(ascending, (tuple, list))
:
```

c:\users\addu\appdata\local\programs\python\python37\lib\si

```

te-packages\pandas\core\generic.py in _get_label_or_level_v
alues(self, key, axis)
    1772         values = self.axes[axis].get_level_valu
es(key)._values
    1773     else:
-> 1774         raise KeyError(key)
    1775
    1776     # Check for duplicates

```

KeyError: 'param_estimator__alpha'

In [36]:


```

results_2 = results_2.sort_values(['param_estimator__alpha'])

train_f1_micro= results_2['mean_train_score']
train_f1_micro_std= results_2['std_train_score']
cv_f1_micro = results_2['mean_test_score']
cv_f1_micro_std= results_2['std_test_score']
K = results_2['param_estimator__alpha']

plt.plot(K, train_f1_micro, label='Train F1 Micro')
plt.plot(K, cv_f1_micro, label='CV F1 Micro')
plt.scatter(K, train_f1_micro, label='Train F1 Micro points')
plt.scatter(K, cv_f1_micro, label='CV F1 Micro points')
plt.legend()
plt.xlabel("alpha : hyperparameter")
plt.xscale('log')
plt.ylabel("F1 Micro")
plt.title("Hyper parameter Vs F1 Micro plot")
plt.grid()
plt.show()

```



3. Try OneVsRestClassifier with Linear-SVM (SGDClassifier with loss-hinge)

In [9]:

```
classifier = OneVsRestClassifier(SGDClassifier(loss = 'hinge', alpha = 0.01))
classifier.fit(X_train_bow, y_train)
predictions = classifier.predict(X_test_bow)

print("Accuracy :", metrics.accuracy_score(y_test, predictions))
print("Hamming loss ", metrics.hamming_loss(y_test, predictions))

precision = precision_score(y_test, predictions, average='micro')
recall = recall_score(y_test, predictions, average='micro')
f1 = f1_score(y_test, predictions, average='micro')

print("Micro-average quality numbers")
print("Precision: {:.4f}, Recall: {:.4f}, F1-measure: {:.4f}".format(precision, recall, f1))

precision = precision_score(y_test, predictions, average='macro')
recall = recall_score(y_test, predictions, average='macro')
f1 = f1_score(y_test, predictions, average='macro')

print("Macro-average quality numbers")
print("Precision: {:.4f}, Recall: {:.4f}, F1-measure: {:.4f}".format(precision, recall, f1))

print(metrics.classification_report(y_test, predictions))
```

Accuracy : 0.20205

Hamming loss 0.00294848

Micro-average quality numbers

Precision: 0.7831, Recall: 0.2100, F1-measure: 0.3312

Macro-average quality numbers

Precision: 0.4435, Recall: 0.1134, F1-measure: 0.1617

	precision	recall	f1-score	support
0	0.96	0.56	0.71	5519
1	0.64	0.09	0.15	8190
2	0.83	0.30	0.44	6529
3	0.79	0.23	0.36	3231
4	0.87	0.35	0.49	6430
5	0.81	0.31	0.45	2879
6	0.88	0.48	0.62	5086
7	0.90	0.50	0.64	4533
8	0.52	0.15	0.23	3000
9	0.83	0.50	0.62	2765
10	0.69	0.05	0.09	3051
11	0.00	0.00	0.00	3000

Conclusion

In [13]:

```
from prettytable import PrettyTable

table = PrettyTable()
table.field_names = ["Vectorizer", "Model", "Hyper Parameters", "Precision",
                    "F1-Micro"]
table.add_row(['TFIDF', 'Logestic Regression (SGD)', 'Alpha = 0.00001', 0.722, 0.49])
table.add_row(['TFIDF', 'Logestic Regression', 'C = 1', 0.7172, 0.4858])
table.add_row(['', '', '', '', ''])
table.add_row(['BOW', 'Logestic Regression', 'C = 0.01', 0.6875, 0.4380])
table.add_row(['BOW', 'Linear SVM (SGD)', 'Alpha = 0.01', 0.7831, 0.3312])
print(table)
```

```
+-----+-----+-----+-----+
+-----+-----+
| Vectorizer | Model | Hyper Parameters | Precision | F1-Micro |
+-----+-----+-----+-----+
+-----+-----+
| TFIDF | Logistic Regression (SGD) | Alpha = 0.00001 | 0.722 | 0.49 |
| TFIDF | Logistic Regression | C = 1 | 0.7172 | 0.4858 |
| | | | |
| | | | |
| BOW | Logistic Regression | C = 0.01 | 0.6875 | 0.438 |
| BOW | Linear SVM (SGD) | Alpha = 0.01 | 0.7831 | 0.3312 |
+-----+-----+-----+-----+
+-----+-----+
```

Summary

- Used 0.5 Million data points and Top 500 Tags
- As the dataset is large, performed 2 folds cross validation using GridSearchCV
- The best model obtained is Linear SVM with Alpha = 0.01
- BOW vectorizer on Linear SVM gave the best Precision of 78.31
- Precision was improved from 72.2 to 78.31